

Data Preparation:

```

import tensorflow as tf
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.preprocessing.image import ImageDataGenerator

(x_train, y_train), (x_test, y_test) = cifar10.load_data()

x_train = x_train.astype("float32") / 255.0
x_test = x_test.astype("float32") / 255.0
y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)

datagen = ImageDataGenerator(
    rotation_range=15,
    width_shift_range=0.1,
    height_shift_range=0.1,
    horizontal_flip=True
)
datagen.fit(x_train)

print("Training images:", x_train.shape)
print("Training labels:", y_train.shape)
print("Testing images:", x_test.shape)
print("Testing labels:", y_test.shape)

```

Training images: (50000, 32, 32, 3)
 Training labels: (50000, 10)
 Testing images: (10000, 32, 32, 3)
 Testing labels: (10000, 10)

Model Implementation:

```

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Flatten, Dense

model = Sequential()
model.add(Flatten(input_shape=(32, 32, 3)))
model.add(Dense(512, activation='relu'))
model.add(Dense(256, activation='relu'))
model.add(Dense(10, activation='softmax'))

/usr/local/lib/python3.12/dist-packages/keras/src/layers/reshaping/flatten.py:37: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential m
super().__init__(**kwargs)

```

```

from tensorflow.keras.optimizers import Adam

model.compile(
    optimizer=Adam(),
    loss='categorical_crossentropy',
    metrics=['accuracy']
)
model.summary()

```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|-------------------|--------------|-----------|
| flatten (Flatten) | (None, 3072) | 0 |
| dense (Dense) | (None, 512) | 1,573,376 |
| dense_1 (Dense) | (None, 256) | 131,328 |
| dense_2 (Dense) | (None, 10) | 2,570 |

Total params: 1,707,274 (6.51 MB)
 Trainable params: 1,707,274 (6.51 MB)
 Non-trainable params: 0 (0.00 B)

```

history = model.fit(
    x_train, y_train,
    epochs=10,
    batch_size=64,
    validation_data=(x_test, y_test)
)
test_loss, test_accuracy = model.evaluate(x_test, y_test)
print("Test Accuracy:", test_accuracy)

```

```

Epoch 1/10
782/782 25s 32ms/step - accuracy: 0.3702 - loss: 1.7476 - val_accuracy: 0.4082 - val_loss: 1.6455
Epoch 2/10
782/782 25s 32ms/step - accuracy: 0.4172 - loss: 1.6300 - val_accuracy: 0.4403 - val_loss: 1.5715
Epoch 3/10
782/782 25s 32ms/step - accuracy: 0.4494 - loss: 1.5473 - val_accuracy: 0.4505 - val_loss: 1.5418
Epoch 4/10
782/782 41s 32ms/step - accuracy: 0.4596 - loss: 1.5140 - val_accuracy: 0.4547 - val_loss: 1.5387
Epoch 5/10
782/782 41s 32ms/step - accuracy: 0.4728 - loss: 1.4758 - val_accuracy: 0.4517 - val_loss: 1.5368
Epoch 6/10
782/782 41s 32ms/step - accuracy: 0.4867 - loss: 1.4427 - val_accuracy: 0.4662 - val_loss: 1.4770
Epoch 7/10
782/782 25s 32ms/step - accuracy: 0.4912 - loss: 1.4270 - val_accuracy: 0.4745 - val_loss: 1.4926
Epoch 8/10
782/782 26s 33ms/step - accuracy: 0.5032 - loss: 1.3886 - val_accuracy: 0.4727 - val_loss: 1.4889
Epoch 9/10
782/782 24s 31ms/step - accuracy: 0.5115 - loss: 1.3674 - val_accuracy: 0.4659 - val_loss: 1.5024
Epoch 10/10
782/782 24s 31ms/step - accuracy: 0.5162 - loss: 1.3455 - val_accuracy: 0.4888 - val_loss: 1.4330
313/313 3s 10ms/step - accuracy: 0.4889 - loss: 1.4310
Test Accuracy: 0.4887999892234802

```

Model Training:

```

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten
from tensorflow.keras.optimizers import Adam
import pandas as pd

model = Sequential([
    Flatten(input_shape=(32, 32, 3)),
    Dense(256, activation='relu'),
    Dense(128, activation='relu'),
    Dense(10, activation='softmax')
])

learning_rate = 0.001
epochs = 20
batch_size = 64

model.compile(
    optimizer=Adam(learning_rate=learning_rate),
    loss='categorical_crossentropy',
    metrics=['accuracy']
)

history = model.fit(
    x_train, y_train,
    epochs=epochs,
    batch_size=batch_size,
    validation_data=(x_test, y_test)
)

history_df = pd.DataFrame(history.history)
history_df.to_csv("ann_training_history.csv", index=False)

model.save("ann_cifar10_model.h5")

Epoch 1/20
782/782 19s 22ms/step - accuracy: 0.2708 - loss: 2.0211 - val_accuracy: 0.3771 - val_loss: 1.7381
Epoch 2/20
782/782 14s 18ms/step - accuracy: 0.3905 - loss: 1.6956 - val_accuracy: 0.4261 - val_loss: 1.6362
Epoch 3/20
782/782 22s 28ms/step - accuracy: 0.4284 - loss: 1.5991 - val_accuracy: 0.4447 - val_loss: 1.5607
Epoch 4/20
782/782 18s 23ms/step - accuracy: 0.4442 - loss: 1.5506 - val_accuracy: 0.4583 - val_loss: 1.5312
Epoch 5/20
782/782 17s 18ms/step - accuracy: 0.4619 - loss: 1.5015 - val_accuracy: 0.4397 - val_loss: 1.5656
Epoch 6/20
782/782 15s 19ms/step - accuracy: 0.4709 - loss: 1.4822 - val_accuracy: 0.4671 - val_loss: 1.5170
Epoch 7/20
782/782 14s 18ms/step - accuracy: 0.4844 - loss: 1.4443 - val_accuracy: 0.4502 - val_loss: 1.5655
Epoch 8/20
782/782 15s 20ms/step - accuracy: 0.4858 - loss: 1.4331 - val_accuracy: 0.4702 - val_loss: 1.5051
Epoch 9/20
782/782 16s 20ms/step - accuracy: 0.4947 - loss: 1.4118 - val_accuracy: 0.4770 - val_loss: 1.4862
Epoch 10/20
782/782 19s 19ms/step - accuracy: 0.5067 - loss: 1.3794 - val_accuracy: 0.4815 - val_loss: 1.4574
Epoch 11/20
782/782 15s 19ms/step - accuracy: 0.5115 - loss: 1.3691 - val_accuracy: 0.4733 - val_loss: 1.4949
Epoch 12/20
782/782 14s 18ms/step - accuracy: 0.5126 - loss: 1.3614 - val_accuracy: 0.4929 - val_loss: 1.4311
Epoch 13/20
782/782 15s 19ms/step - accuracy: 0.5251 - loss: 1.3294 - val_accuracy: 0.4903 - val_loss: 1.4448
Epoch 14/20
782/782 14s 18ms/step - accuracy: 0.5283 - loss: 1.3220 - val_accuracy: 0.4929 - val_loss: 1.4253
Epoch 15/20
782/782 14s 18ms/step - accuracy: 0.5324 - loss: 1.3106 - val_accuracy: 0.4917 - val_loss: 1.4286
Epoch 16/20
782/782 15s 19ms/step - accuracy: 0.5399 - loss: 1.2910 - val_accuracy: 0.5030 - val_loss: 1.4090
Epoch 17/20
782/782 21s 20ms/step - accuracy: 0.5396 - loss: 1.2890 - val_accuracy: 0.4914 - val_loss: 1.4351
Epoch 18/20
782/782 14s 18ms/step - accuracy: 0.5466 - loss: 1.2675 - val_accuracy: 0.5033 - val_loss: 1.4050
Epoch 19/20
782/782 15s 19ms/step - accuracy: 0.5553 - loss: 1.2508 - val_accuracy: 0.4974 - val_loss: 1.4146
Epoch 20/20
782/782 14s 17ms/step - accuracy: 0.5513 - loss: 1.2544 - val_accuracy: 0.5071 - val_loss: 1.4080
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the

```

Model Evaluation:

```

import numpy as np
from sklearn.metrics import confusion_matrix

y_pred = model.predict(x_test)
y_pred_classes = np.argmax(y_pred, axis=1)
y_true = np.argmax(y_test, axis=1)

cm = confusion_matrix(y_true, y_pred_classes)
print("Confusion Matrix:\n", cm)
test_loss, test_accuracy = model.evaluate(x_test, y_test, verbose=0)

print("Test Loss:", test_loss)
print("Test Accuracy:", test_accuracy)

```

```

313/313 4s 11ms/step
Confusion Matrix:
[[534 19 116 15 49 21 30 42 125 49]
 [ 47 552 33 21 8 27 18 41 72 181]
 [ 47 13 420 56 111 89 135 90 21 18]
 [ 22 14 123 194 63 256 161 94 30 43]
 [ 34 3 141 27 407 53 161 138 22 14]
 [ 16 9 96 106 57 433 114 126 22 21]
 [ 3 5 73 50 85 48 656 55 9 16]
 [ 26 2 57 25 79 58 38 658 12 45]
 [114 34 31 14 35 24 16 26 617 89]
 [ 44 125 19 25 17 21 39 68 42 600]]
Test Loss: 1.4080229997634888
Test Accuracy: 0.507099986076355

```

Performance Comparison:

```

import matplotlib.pyplot as plt
plt.figure()

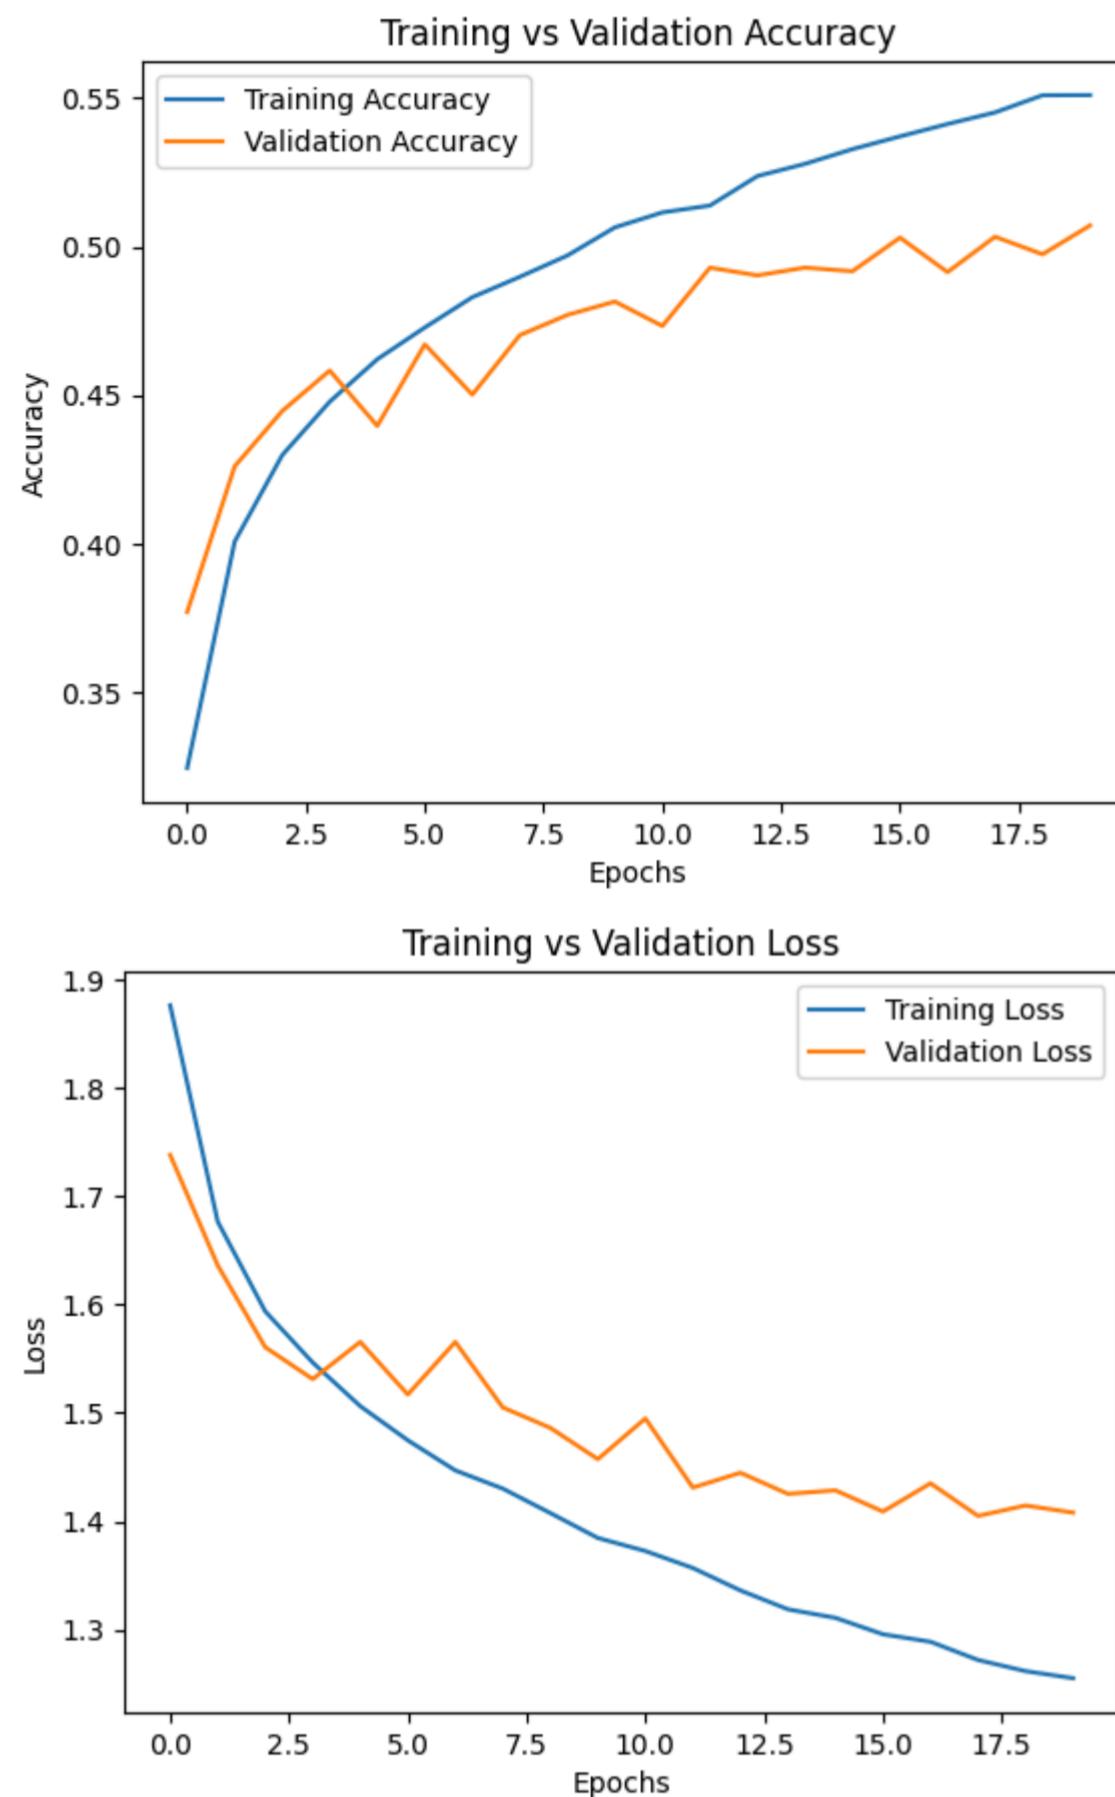
```

```

plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.title('Training vs Validation Accuracy')
plt.legend()
plt.show()

plt.figure()
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Training vs Validation Loss')
plt.legend()
plt.show()

```



Summary:

The CIFAR-10 dataset was loaded to perform image classification using an Artificial Neural Network (ANN). The dataset consists of 60,000 color images of size $32 \times 32 \times 3$ belonging to 10 classes, with 50,000 images used for training and 10,000 images for testing. Data preprocessing involved normalizing pixel values to the range [0–1] and one-hot encoding the class labels. To improve generalization, data augmentation techniques such as rotation, width and height shifting, and horizontal flipping were applied.

An ANN model was developed using a Flatten layer followed by fully connected Dense layers with ReLU activation and a Softmax output layer. The model was trained using the Adam optimizer and categorical cross-entropy loss with a batch size of 64 for 10 and 20 epochs as shown in the experiments. The trained model achieved an approximate training accuracy of 55% and a testing accuracy of about 50%.

The results indicate that model architecture, batch size, and number of epochs play a significant role in determining classification performance.